The aim is to build a full-duplex, voice-enabled AI agent running on a Raspberry Pi (4 or 5). Your responsibilities:

- Install AI agent harness (e.g. onit, pipecat, crew) with MCP support on RPi
- Build MCP servers that to SOTA ASR and TTS
- Install VAD on RPi for voice detection,
- Run a full-duplex AI agent that can real-world tasks
- You may use commercial APIs for models but you build and run your MCP servers (on RPi)
- Build a latency benchmark (classwork): a standardized latency measurement method. Note that this is . Speech interaction is natural, conversational and interruptible.

---

sabi ni sir, main focus daw is yung same top 10 most common tasks from ME2 then top up nalang ng more complex tasks. If tama tong memory ko, eh magfocus muna dun sa 10 tasks na yun before natin lagyan ng mga more complex tasks

Sbi ni sir, whatever onit can do, the vcm should also be able to do

Q&A kanina
- Can we use DGX to setup 24/7 inference of raw ASR/TTS/etc models? YES raw. Avoid hosting LLMs raw since may QWEN.

Note: Coordinate tayo sa sharc if they can setup an inference or api distribution for us
- Can we build ontop of Pipecat? Yes raw (i think this makes the ME more bearable)